

****EDA Vehicle Insurance project****

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb

data = pd.read_csv("C:\\\\Users\\\\Nikhil\\\\Downloads\\\\
Vehicle_Insurance.csv")
data
```

	id	Gender	Age	Driving_License	Region_Code	
Previously_Insured \						
0	1	Male	44	1	28.0	
0	2	Male	76	1	3.0	
0	3	Male	47	1	28.0	
0	4	Male	21	1	11.0	
1	5	Female	29	1	41.0	
1	
...	
381104	381105	Male	74	1	26.0	
1	381105	381106	Male	30	1	37.0
1	381106	381107	Male	21	1	30.0
1	381107	381108	Female	68	1	14.0
0	381108	381109	Male	46	1	29.0
0	
	Vehicle_Age	Vehicle_Damage	Annual_Premium			
Policy_Sales_Channel \						
0	> 2 Years	Yes	40454.0			
26.0						
1	1-2 Year	No	33536.0			
26.0						
2	> 2 Years	Yes	38294.0			
26.0						
3	< 1 Year	No	28619.0			
152.0						
4	< 1 Year	No	27496.0			
152.0						
...	

```

.
381104    1-2 Year           No      30170.0
26.0
381105    < 1 Year          No      40016.0
152.0
381106    < 1 Year          No      35118.0
160.0
381107    > 2 Years         Yes     44617.0
124.0
381108    1-2 Year          No      41777.0
26.0

```

	Vintage	Response
0	217	1
1	183	0
2	27	1
3	203	0
4	39	0
..
381104	88	0
381105	131	0
381106	161	0
381107	74	0
381108	237	0

[381109 rows x 12 columns]

data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 381109 entries, 0 to 381108
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                381109 non-null   int64  
 1   Gender             381109 non-null   object 
 2   Age                381109 non-null   int64  
 3   Driving_License    381109 non-null   int64  
 4   Region_Code        381109 non-null   float64
 5   Previously_Insured 381109 non-null   int64  
 6   Vehicle_Age        381109 non-null   object 
 7   Vehicle_Damage     381109 non-null   object 
 8   Annual_Premium     381109 non-null   float64
 9   Policy_Sales_Channel 381109 non-null   float64
 10  Vintage            381109 non-null   int64  
 11  Response           381109 non-null   int64  
dtypes: float64(3), int64(6), object(3)
memory usage: 34.9+ MB

```

Data Cleaning

```
data.isna().sum()

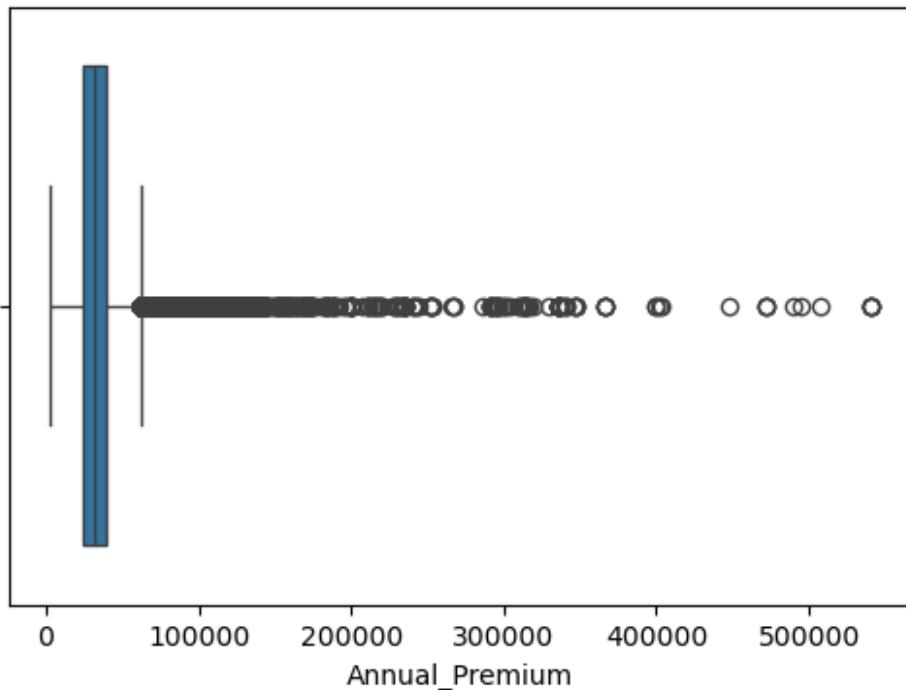
id          0
Gender      0
Age         0
Driving_License 0
Region_Code 0
Previously_Insured 0
Vehicle_Age   0
Vehicle_Damage 0
Annual_Premium 0
Policy_Sales_Channel 0
Vintage      0
Response     0
dtype: int64
```

insight data is clean already as it has no null value

```
data.duplicated().sum()
np.int64(0)

#outlier detection
plt.figure(figsize=(6,4))
sb.boxplot(x=data["Annual_Premium"])
plt.title("Outlier Detection - Annual Premium")
plt.show()
```

Outlier Detection - Annual Premium



```

Q1 = data["Annual_Premium"].quantile(0.25)
Q3 = data["Annual_Premium"].quantile(0.75)
IQR = Q3 - Q1

lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR

data = data[(data["Annual_Premium"] >= lower) &
            (data["Annual_Premium"] <= upper)]

data.info()

<class 'pandas.core.frame.DataFrame'>
Index: 370789 entries, 0 to 381108
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               370789 non-null   int64  
 1   Gender            370789 non-null   object  
 2   Age                370789 non-null   int64  
 3   Driving_License    370789 non-null   int64  
 4   Region_Code        370789 non-null   float64 
 5   Previously_Insured 370789 non-null   int64  
 6   Vehicle_Age        370789 non-null   object  
 7   Vehicle_Damage      370789 non-null   object  
 8   Annual_Premium      370789 non-null   float64 
 9   Policy_Sales_Channel 370789 non-null   float64 
 10  Vintage             370789 non-null   int64  
 11  Response            370789 non-null   int64  
dtypes: float64(3), int64(6), object(3)
memory usage: 36.8+ MB

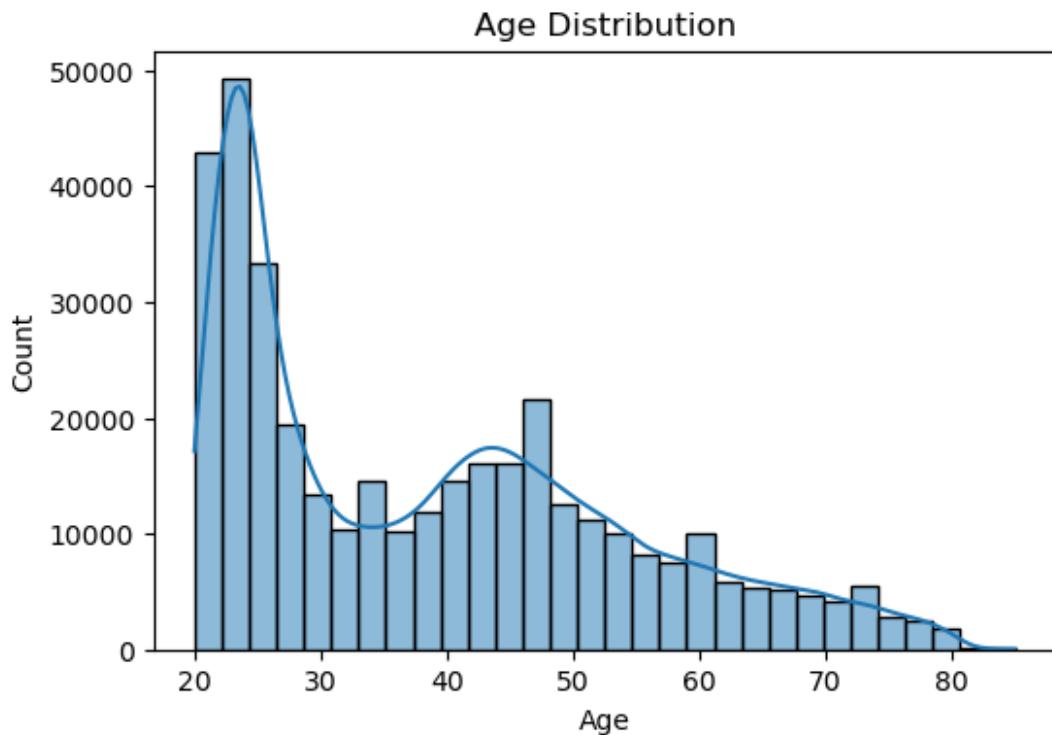
```

Data Visualization

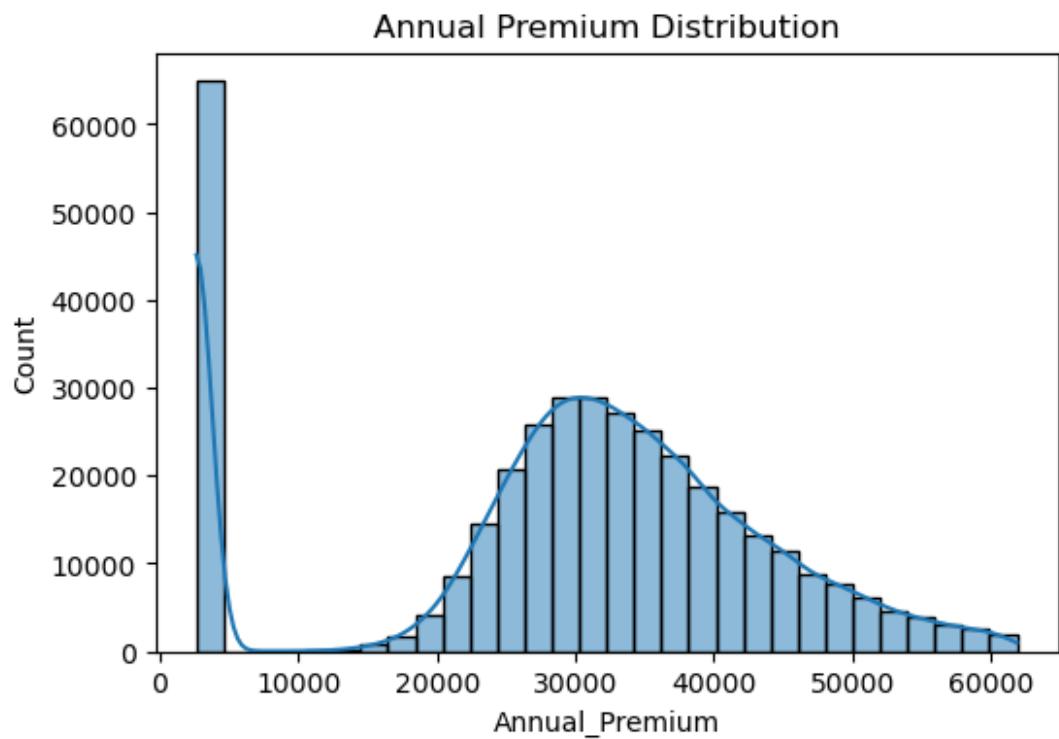
```

#age distribution
plt.figure(figsize=(6,4))
sb.histplot(data["Age"], bins=30, kde=True)
plt.title("Age Distribution")
plt.show()

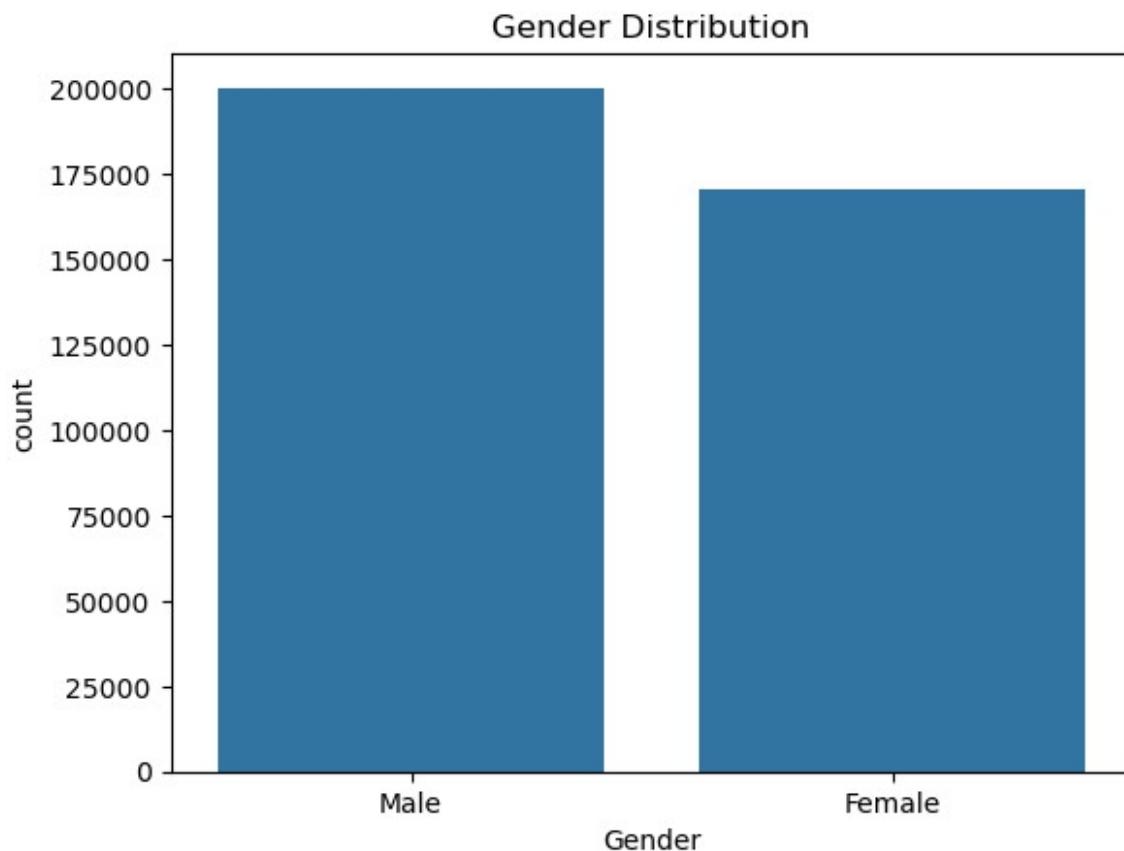
```



```
#annual premium distribution
plt.figure(figsize=(6,4))
sb.histplot(data["Annual_Premium"], bins=30, kde=True)
plt.title("Annual Premium Distribution")
plt.show()
```



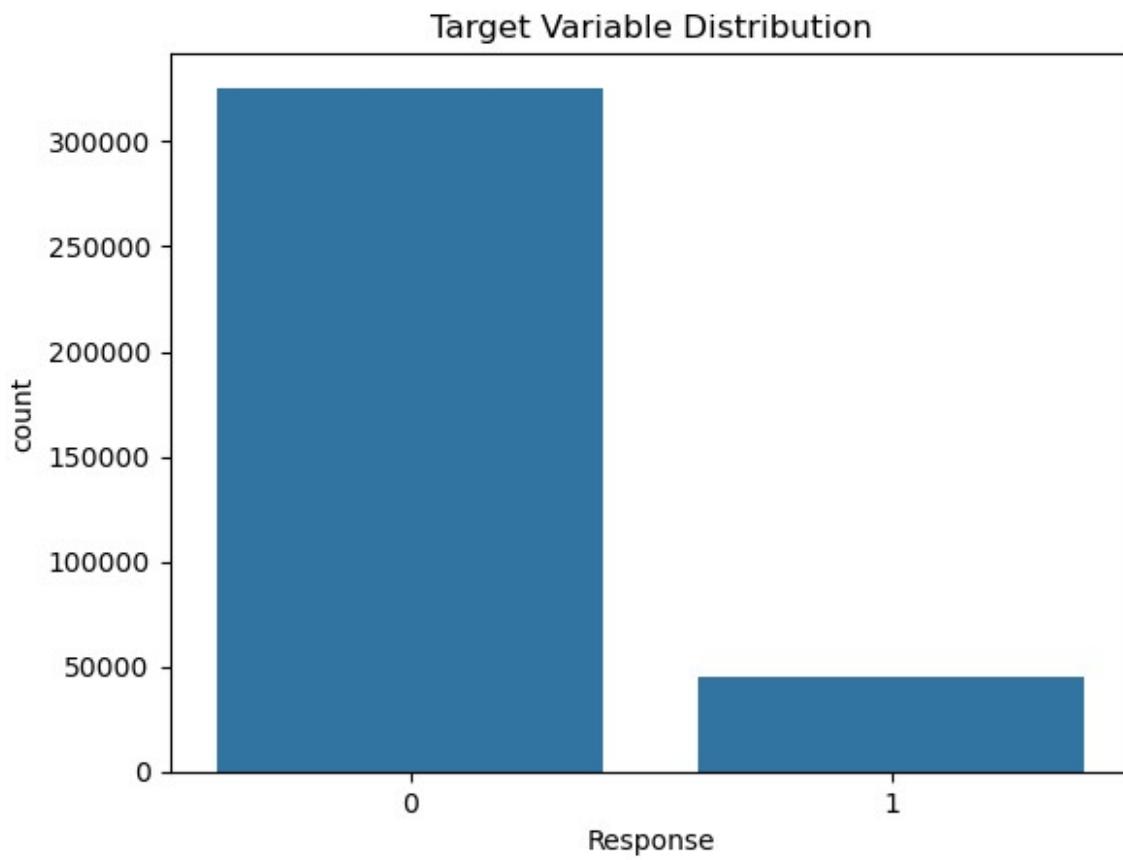
```
#gender distribution  
sb.countplot(x="Gender", data=data)  
plt.title("Gender Distribution")  
plt.show()
```



FEATURE ANALYSIS

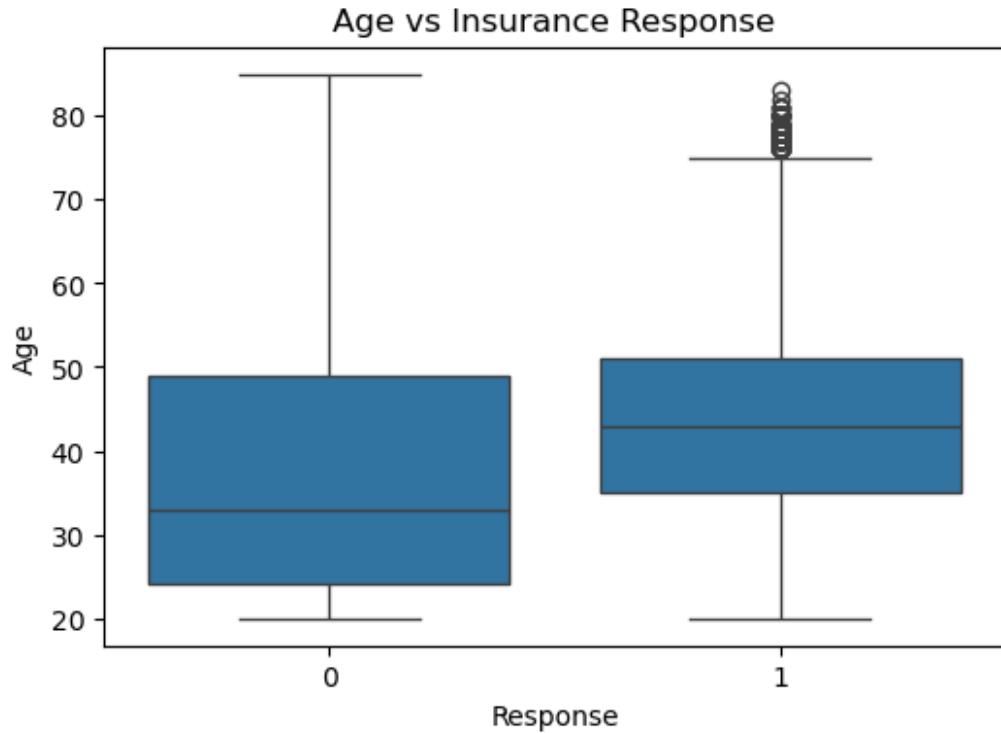
(Relationship between features & target variable – Response)

```
#target Variable Distribution  
sb.countplot(x="Response", data=data)  
plt.title("Target Variable Distribution")  
plt.show()
```



AGE DISTRIBUTION

```
#age vs response
plt.figure(figsize=(6,4))
sb.boxplot(x="Response", y="Age", data=data)
plt.title("Age vs Insurance Response")
plt.show()
```



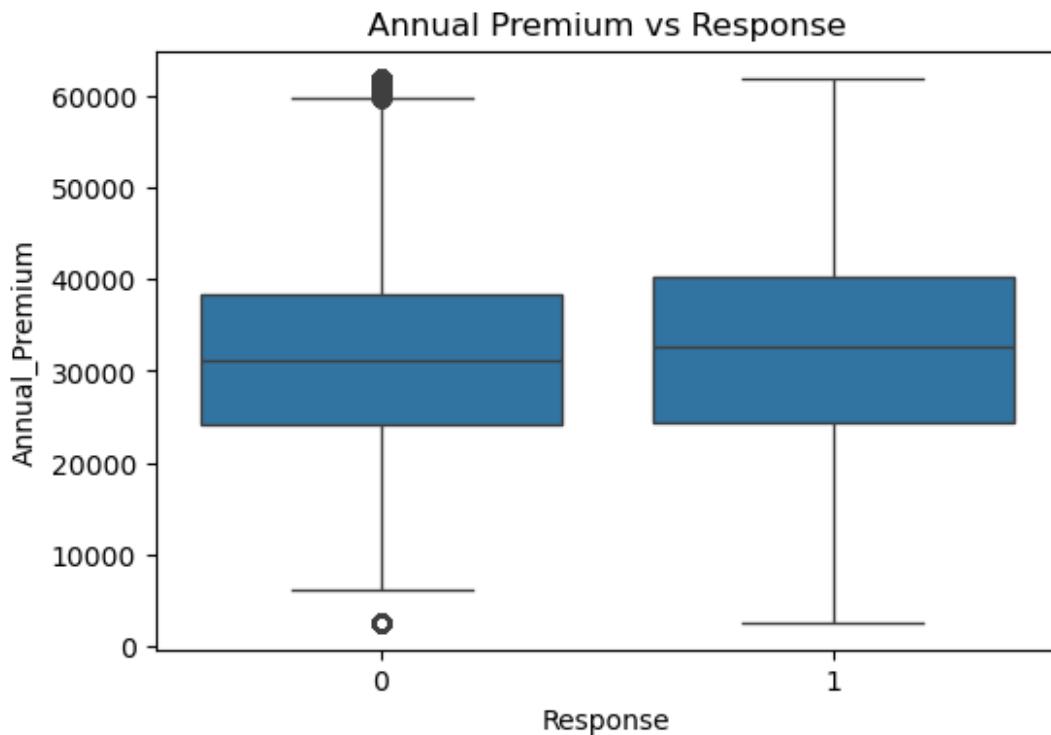
```
data.groupby("Response")["Age"].mean()
```

```
Response
0    38.032494
1    43.270180
Name: Age, dtype: float64
```

Insight: Older customers are more likely to respond to insurance offers.

PREMIUM ANALYSIS

```
#premium vs response
plt.figure(figsize=(6,4))
sb.boxplot(x="Response", y="Annual_Premium", data=data)
plt.title("Annual Premium vs Response")
plt.show()
```

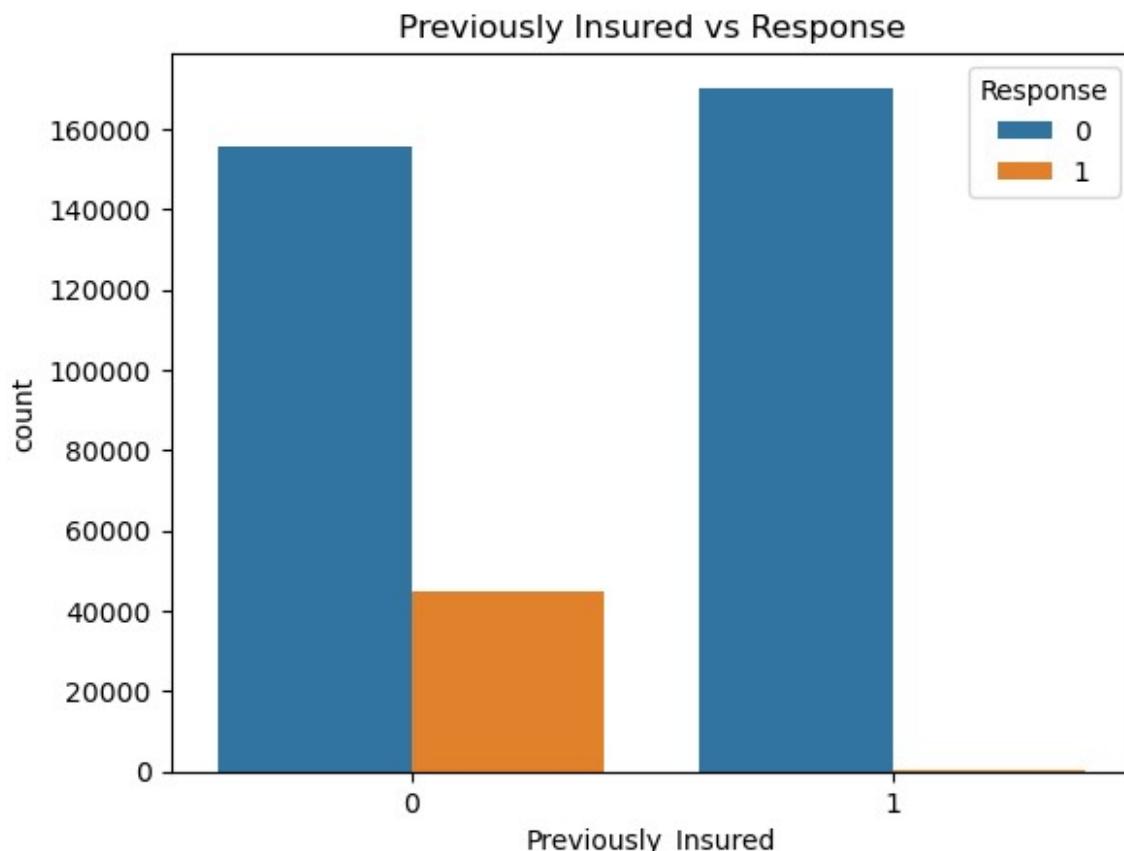


```
data.groupby("Response")["Annual_Premium"].mean()
```

```
Response
0    29162.717548
1    29999.683490
Name: Annual_Premium, dtype: float64
```

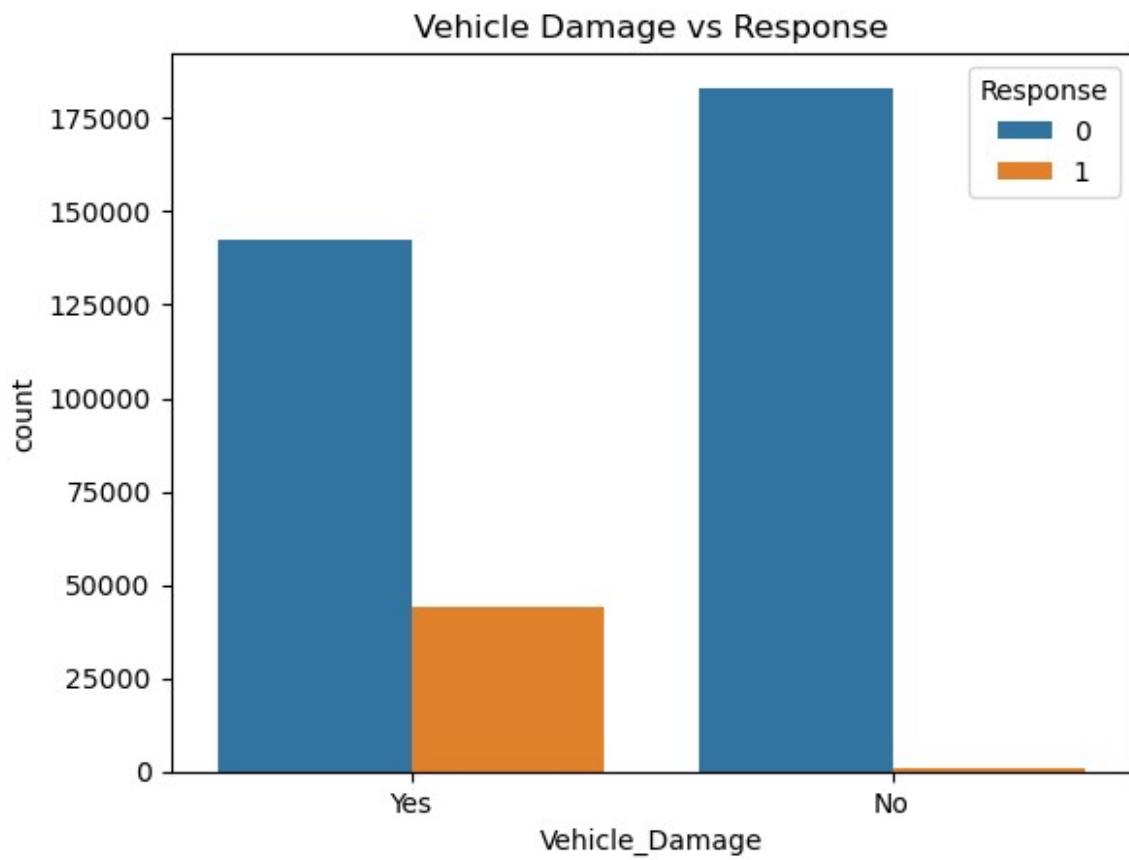
CLAIM FREQUENCIES

```
#previously insured vs response
sb.countplot(x="Previously_Insured", hue="Response", data=data)
plt.title("Previously Insured vs Response")
plt.show()
```



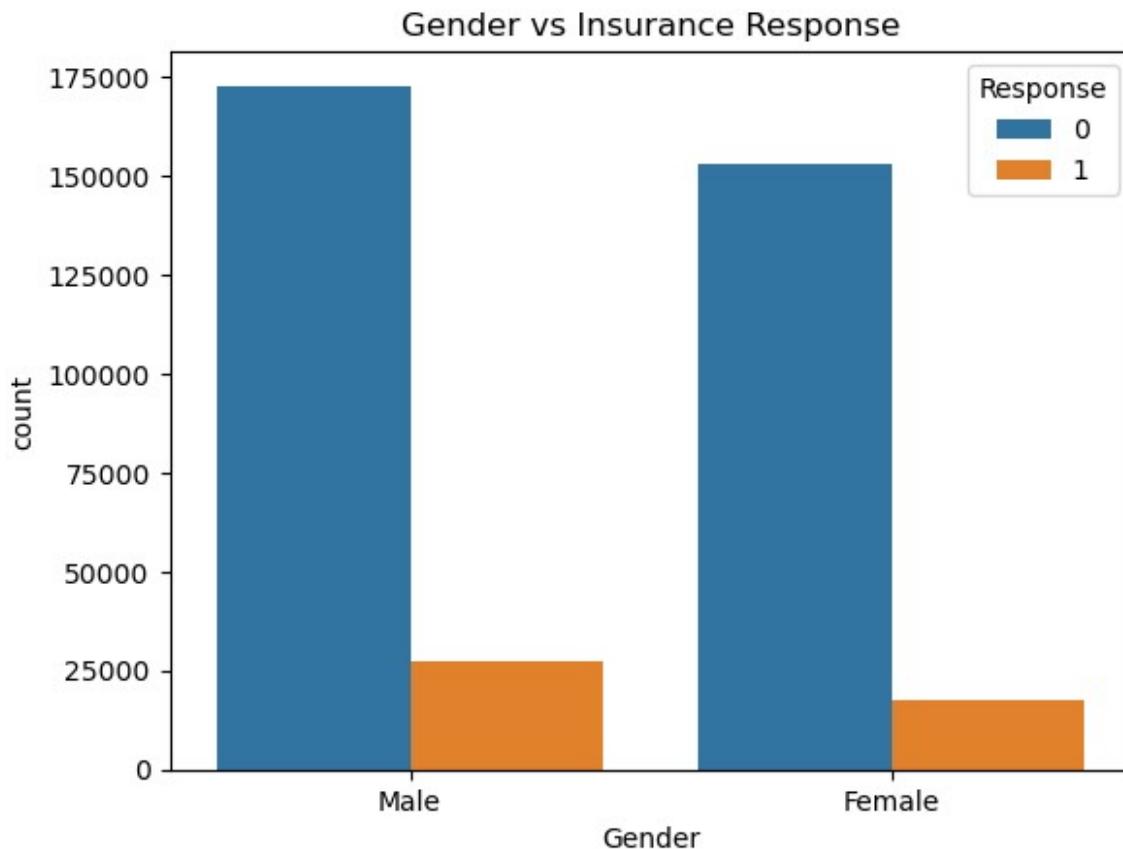
Insight: Customers not previously insured are far more likely to respond.

```
#Vehicle Damage vs Response  
sb.countplot(x="Vehicle_Damage", hue="Response", data=data)  
plt.title("Vehicle Damage vs Response")  
plt.show()
```



Gender analysis

```
#Role of gender in insurance claims
sb.countplot(x="Gender", hue="Response", data=data)
plt.title("Gender vs Insurance Response")
plt.show()
```



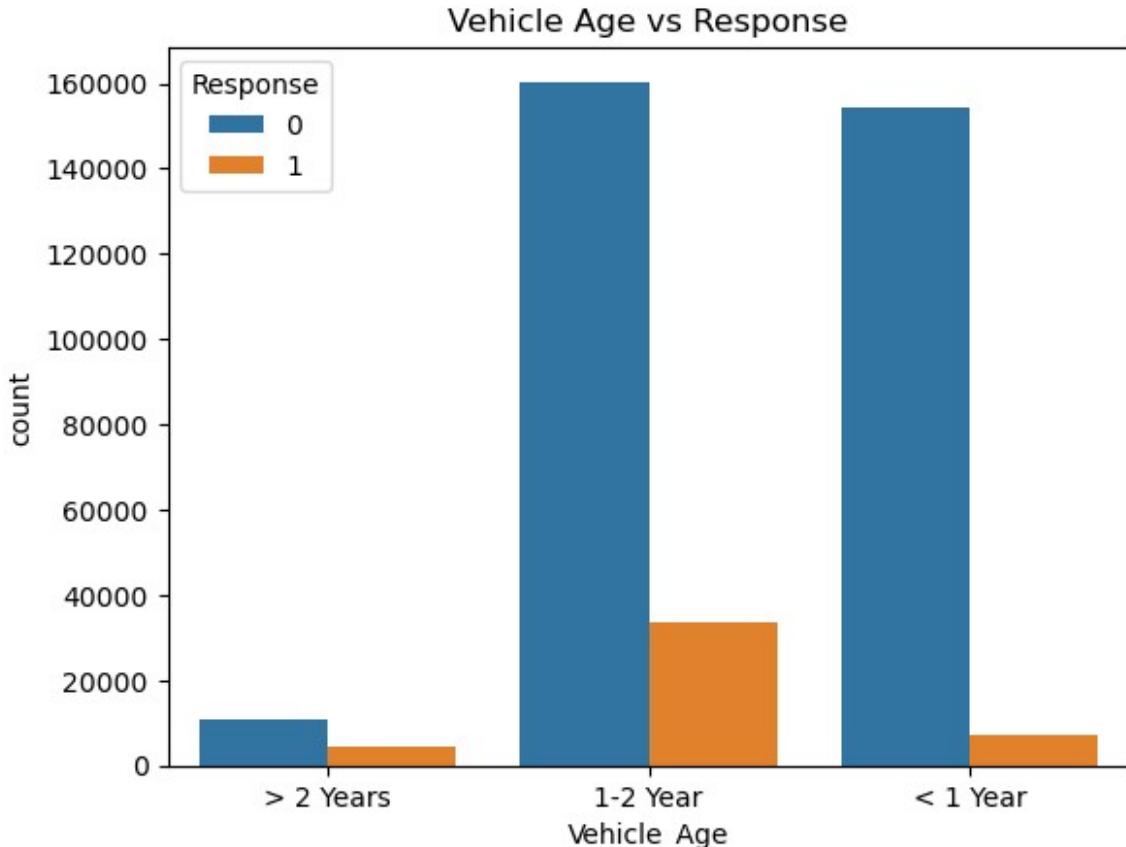
```
data.groupby("Gender")["Response"].mean()

Gender
Female    0.103238
Male      0.137561
Name: Response, dtype: float64
```

Insight: Males show slightly higher response rates than females.

VEHICLE AGE & CLAIMS

```
#Vehicle Age vs Response
sb.countplot(x="Vehicle_Age", hue="Response", data=data)
plt.title("Vehicle Age vs Response")
plt.show()
```

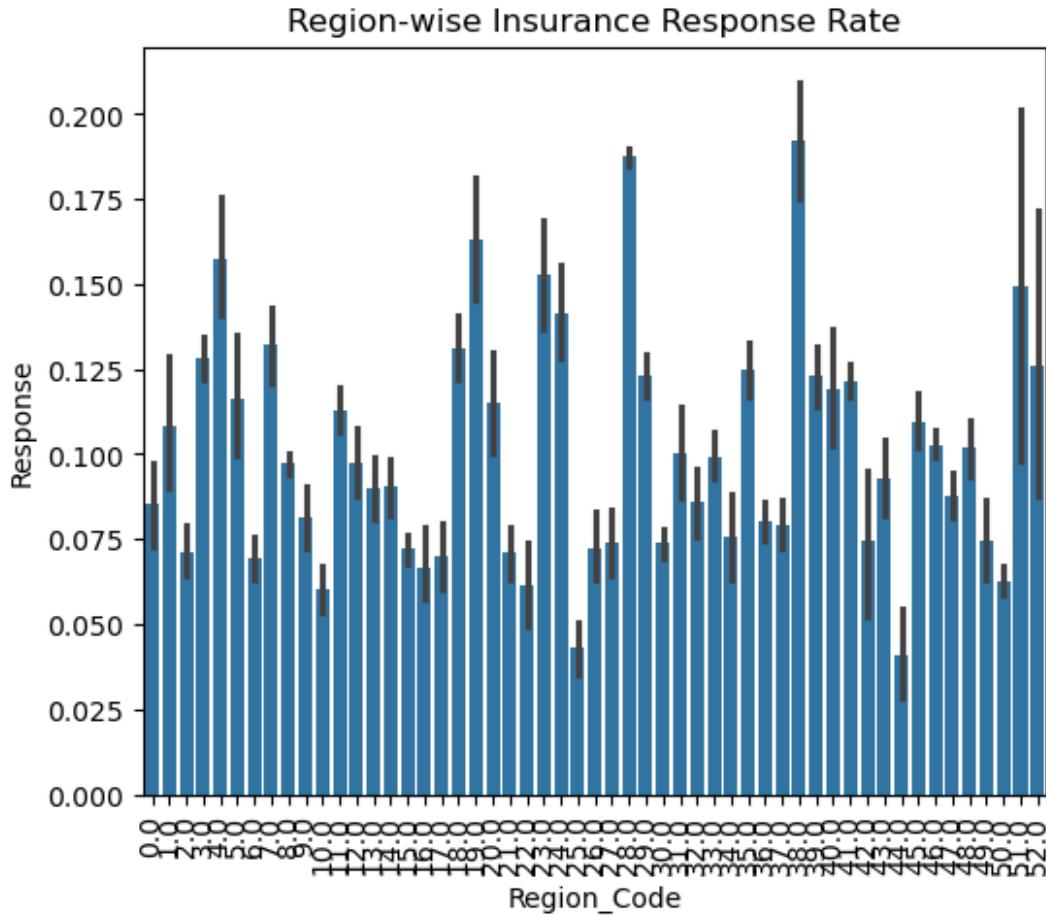


```
data.groupby("Vehicle_Age")["Response"].mean()

Vehicle_Age
1-2 Year      0.173753
< 1 Year      0.043702
> 2 Years     0.289421
Name: Response, dtype: float64
```

REGION-WISE ANALYSIS (Analyze regional patterns in insurance claims)

```
#Region vs Response
plt.figure(figsize=(6,5))
sb.barplot(
    x="Region_Code",
    y="Response",
    data=data,
    estimator=np.mean
)
plt.title("Region-wise Insurance Response Rate")
plt.xticks(rotation=90)
plt.show()
```



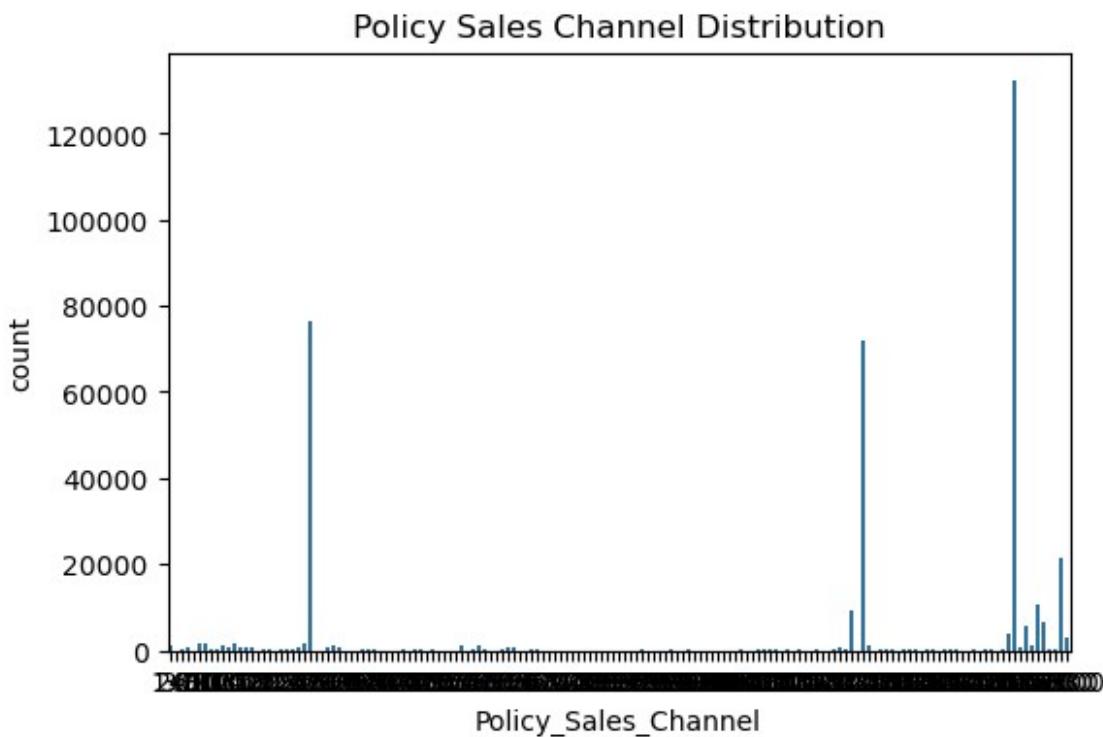
```
data.groupby("Region_Code")
[ "Response"].mean().sort_values(ascending=False).head(10)
```

```
Region_Code
38.0    0.192423
28.0    0.187526
19.0    0.162973
4.0     0.157572
23.0    0.152707
51.0    0.149425
24.0    0.141611
7.0     0.132260
18.0    0.130987
3.0     0.128325
Name: Response, dtype: float64
```

Insight: Certain regions show significantly higher response rates, useful for targeted marketing.

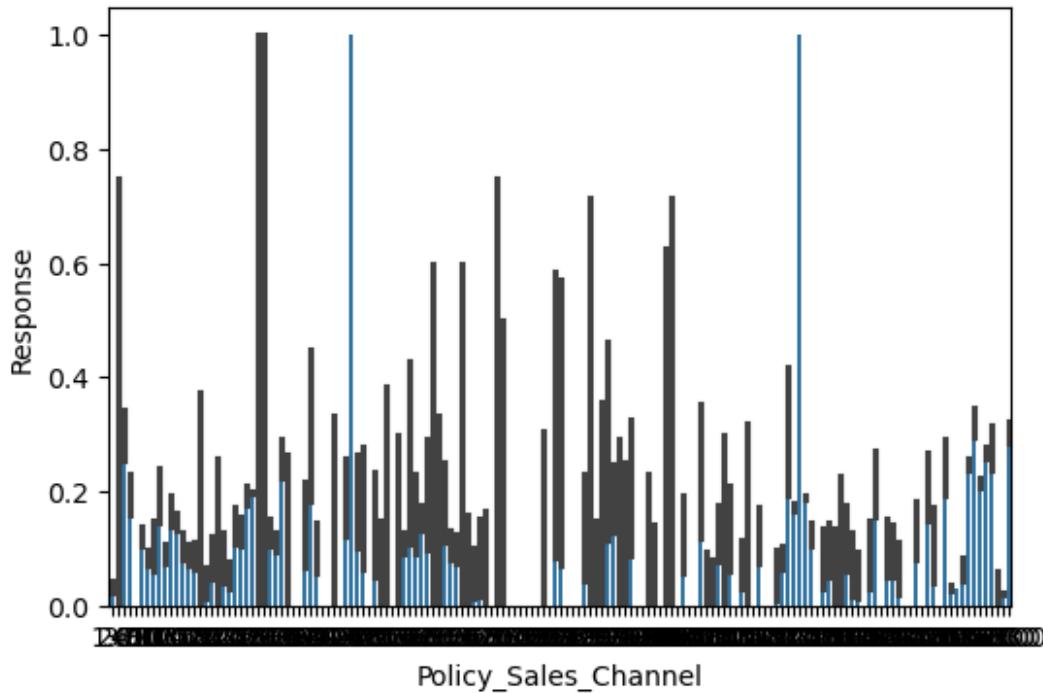
POLICY ANALYSIS (Distribution & impact of different policy types)

```
#Policy Sales Channel Distribution
plt.figure(figsize=(6,4))
sb.countplot(x="Policy_Sales_Channel", data=data)
plt.title("Policy Sales Channel Distribution")
plt.show()
```



```
#Policy Sales Channel vs Response
plt.figure(figsize=(6,4))
sb.barplot(
    x="Policy_Sales_Channel",
    y="Response",
    data=data,
    estimator=np.mean
)
plt.title("Policy Sales Channel vs Response Rate")
plt.show()
```

Policy Sales Channel vs Response Rate



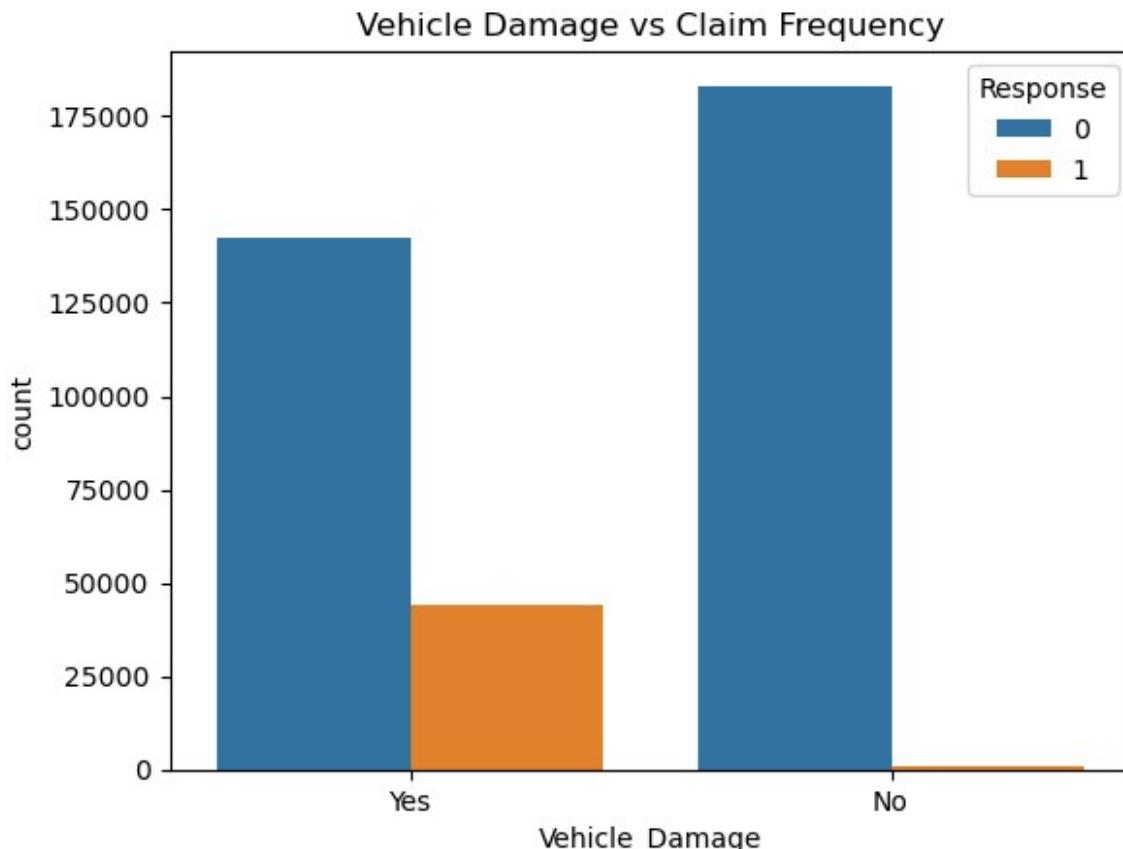
```
data.groupby("Policy_Sales_Channel")["Response"].mean()

Policy_Sales_Channel
1.0      0.032588
2.0      0.250000
3.0      0.298354
4.0      0.193939
6.0      0.000000
...
157.0    0.268072
158.0    0.277207
159.0    0.019608
160.0    0.021918
163.0    0.303571
Name: Response, Length: 155, dtype: float64
```

Insight: Certain sales channels outperform others in customer conversions.

CLAIM FREQUENCY BY VEHICLE DAMAGE

```
sb.countplot(
    x="Vehicle_Damage",
    hue="Response",
    data=data
)
plt.title("Vehicle Damage vs Claim Frequency")
plt.show()
```



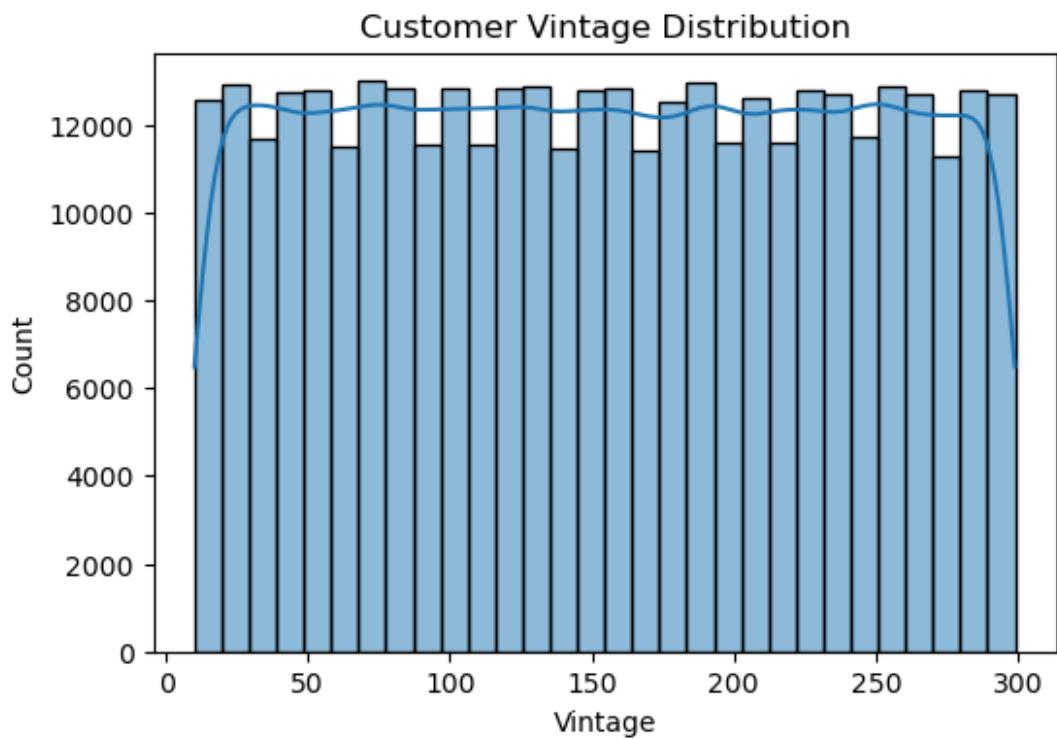
```
data.groupby("Vehicle_Damage")["Response"].mean()

Vehicle_Damage
No      0.005249
Yes     0.236856
Name: Response, dtype: float64
```

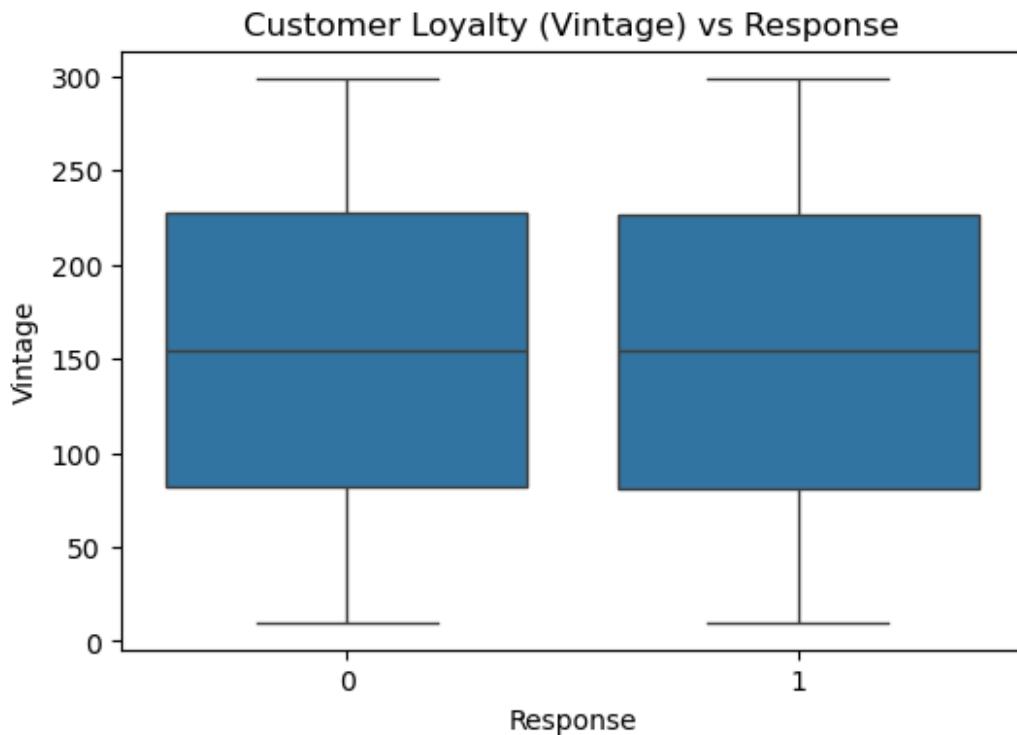
Insight: Customers with vehicle damage = Yes have a much higher probability of filing claims / responding.

CUSTOMER LOYALTY ANALYSIS

```
#Vintage Distribution
plt.figure(figsize=(6,4))
sb.histplot(data["Vintage"], bins=30, kde=True)
plt.title("Customer Vintage Distribution")
plt.show()
```



```
#Vintage vs Response
plt.figure(figsize=(6,4))
sb.boxplot(x="Response", y="Vintage", data=data)
plt.title("Customer Loyalty (Vintage) vs Response")
plt.show()
```



```
data.groupby("Response")["Vintage"].mean()

Response
0    154.396261
1    153.978961
Name: Vintage, dtype: float64
```

Insight: Long-term customers show slightly higher trust, but newer customers respond more aggressively to offers.

TIME ANALYSIS (Temporal patterns in insurance claims)

```
#Vintage Bins (Customer Age Buckets)
data["Vintage_Group"] = pd.cut(
    data["Vintage"],
    bins=[0, 100, 200, 300],
    labels=["New Customers", "Mid-term Customers", "Long-term
Customers"]
)
```

```
C:\Users\Nikhil\AppData\Local\Temp\ipykernel_7252\2520873319.py:2:
SettingWithCopyWarning:
```

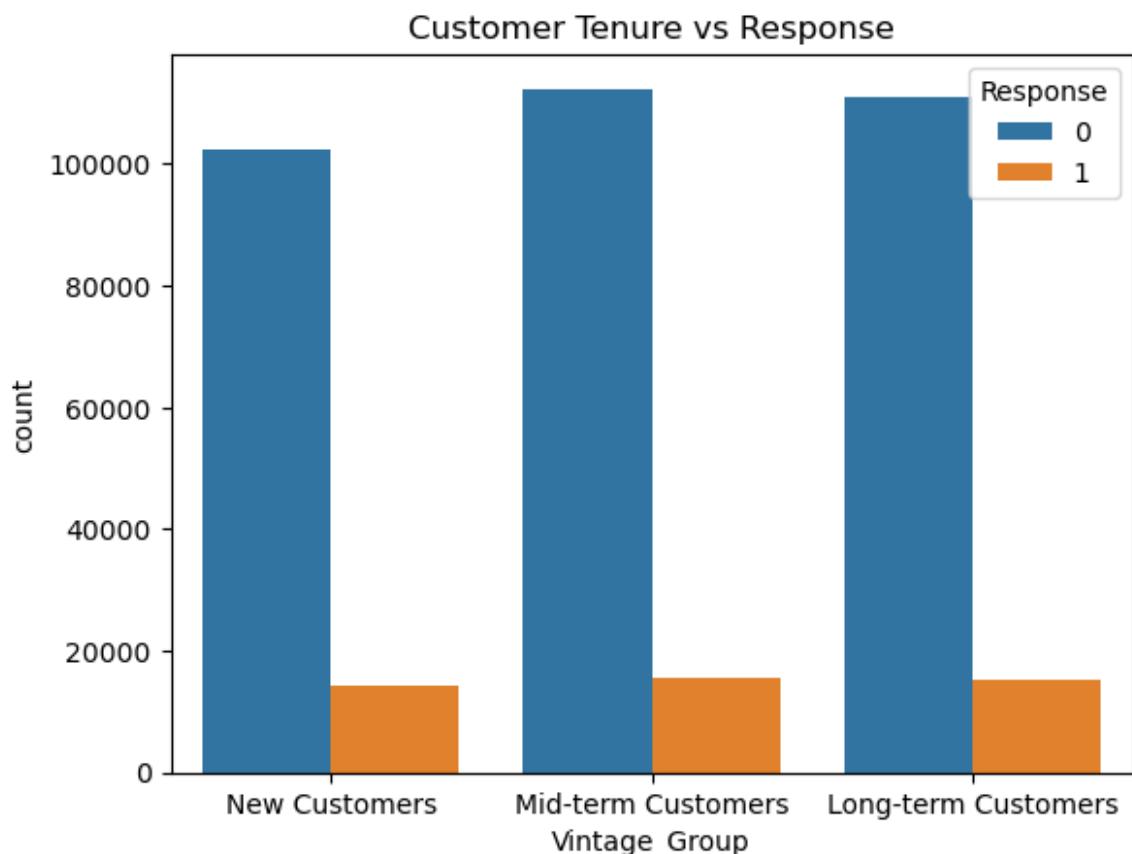
```
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#

```

returning-a-view-versus-a-copy
data["Vintage_Group"] = pd.cut(
    #Vintage Group vs Response
    sb.countplot(
        x="Vintage_Group",
        hue="Response",
        data=data
    )
    plt.title("Customer Tenure vs Response")
    plt.show()
)

```



```

data.groupby("Vintage_Group")["Response"].mean()

C:\Users\Nikhil\AppData\Local\Temp\ipykernel_7252\2877602394.py:1:
FutureWarning: The default of observed=False is deprecated and will be
changed to True in a future version of pandas. Pass observed=False to
retain current behavior or observed=True to adopt the future default
and silence this warning.
    data.groupby("Vintage_Group")["Response"].mean()

Vintage_Group
New Customers          0.122075

```

```
Mid-term Customers      0.121727
Long-term Customers     0.121564
Name: Response, dtype: float64
```

Insight: Customers in the mid-tenure phase show the highest response rate.

Key summary:

Vehicle damage is the strongest predictor of claims

Previously uninsured customers are high-value targets

Middle-aged and senior customers respond more

Certain regions and policy channels outperform others

Mid-tenure customers show higher engagement

Business Recommendations:

Focus marketing on damaged-vehicle owners

Target uninsured customers

Optimize high-performing sales channels

Region-specific campaigns can improve conversion